

Memory & I/Os

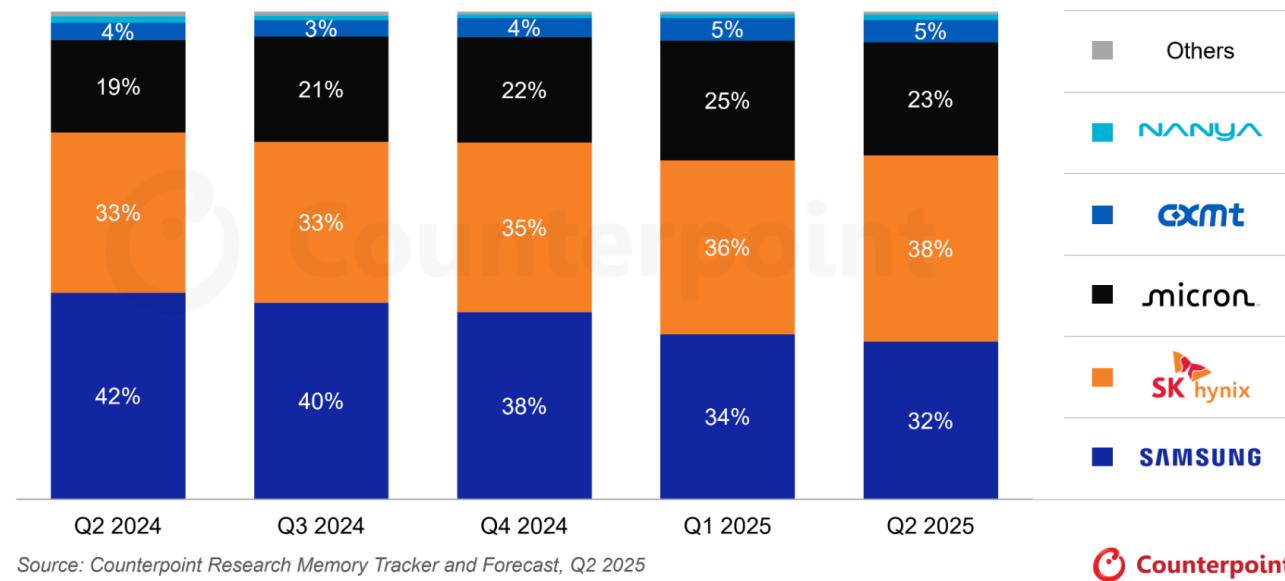
Philippe Bonnet, bonnet@di.ku.dk
HPPS 2025 – 2b

Outline

- Memory
- Pointers
- Byte ordering
- Serialization

DRAM market

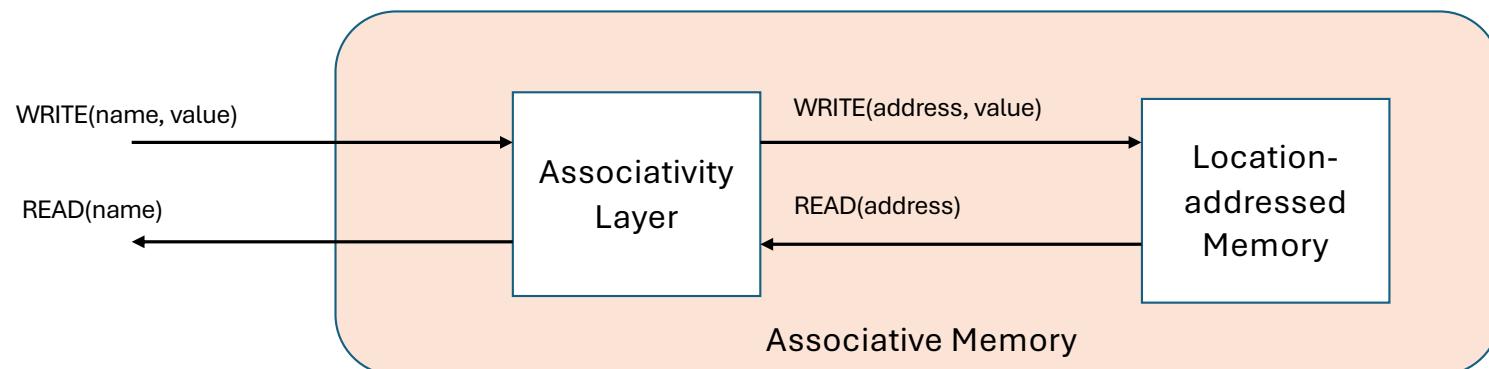
Global DRAM Market Share by Revenue (Q2 2024 - Q2 2025)



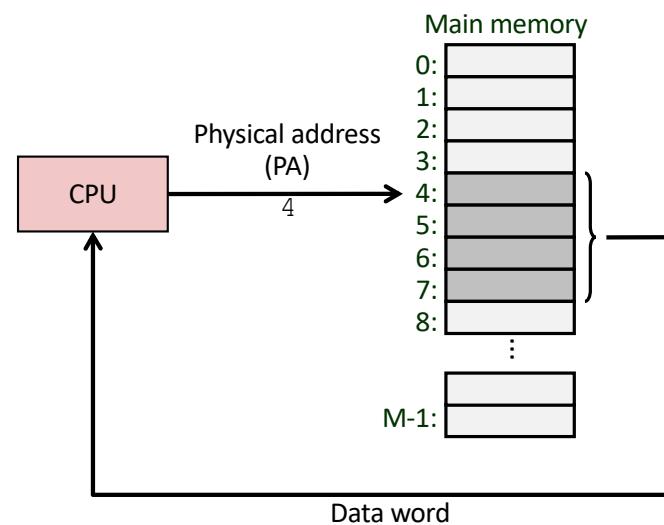
DRAM includes
HBM, DDR5, DDR4

Source: Saltzer and Kaashoek

Memory Abstraction



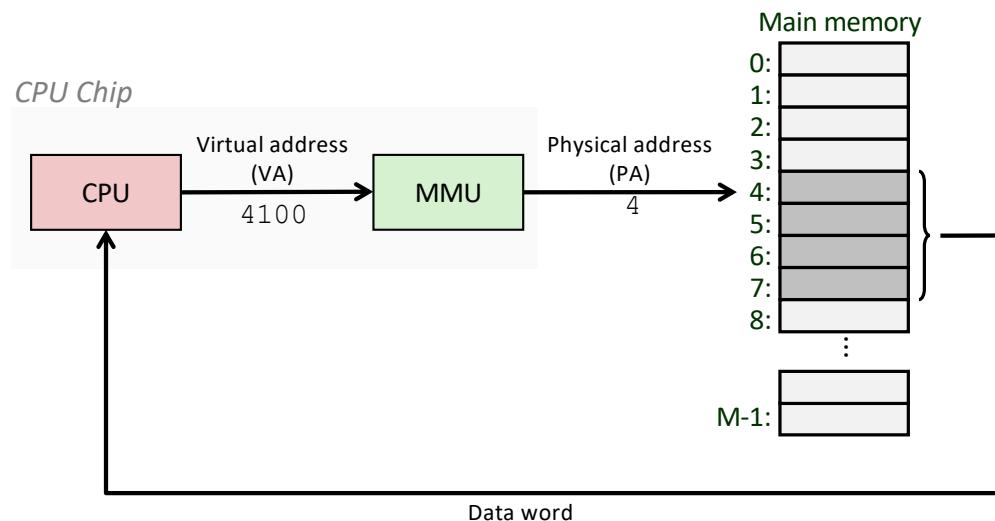
Physical addressing



Used in “simple” embedded systems, e.g., kitchen appliances

Virtual addressing

MMU: Memory Management Unit



- Used in all modern servers, desktops, and laptops
- One of the great ideas in computer science

Address space

- An address space is an ordered set of contiguous addresses (non-negative integers).
- Each word has an address . Address points to 1st byte of word.
On 64 bits machine, each word is 8B.
- Physical address space => associated to RAM
- Virtual address space => associated to each process
(see Troels intro and more later)

Pointers

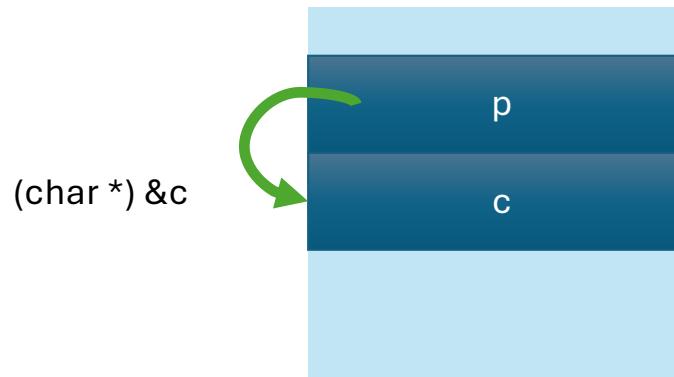
- "A pointer is a variable that contains the address of a variable."
K & R
- We can access the address of any variable in a C program
 - `&x` gives us the (virtual) address of `x`
 - If `x` has type `T`, `&x` has type `T*`

Pointers

p is a char pointer.

It is a variable that contains the address of a char variable.

```
char *p;  
char c;  
p = &c;
```



Why pointers in C?

- To share data without copies
- To manage indirections (e.g., linked list)
- To manage data placement / locality and memory allocation

Beware initialization

```
1 #include <stdio.h>
2
3 int main(void) {
4     int *ptr;
5     *ptr = 20;
6     printf("%d\n", *ptr);
7     return 0;
8 }
```

```
htl719@kumac ~/D/C/H/h/week2> gcc -Wall -Werror ex.c -o ex
ex.c:5:4: error: variable 'ptr' is uninitialized when used here [-Werror,-Wuninitialized]
  5 |     *ptr = 20;
  |     ^
ex.c:4:11: note: initialize the variable 'ptr' to silence this warning
  4 |     int *ptr;
  |             ^
  |             = NULL
1 error generated.
```

Pointers

- Pointers are valid, null or indeterminate.
- A pointer is null when assigned 0
- Null pointers evaluate to false in logical expressions
- Dereferencing indeterminate pointers leads to undefined behaviour

Always initialize pointers!

Byte ordering

Consider a 64 bit integer, i.e., 8 bytes.

How are these 8 bytes ordered in memory?

- **Big endian**

- *Least significant byte has highest address (“comes last”).*
- *Network: TCP*

- **Little endian**

- *Least significant byte has highest address (“comes first”).*
- x86, ARM, Risc-V CPUs

Byte re-ordering when sending/receiving on the network

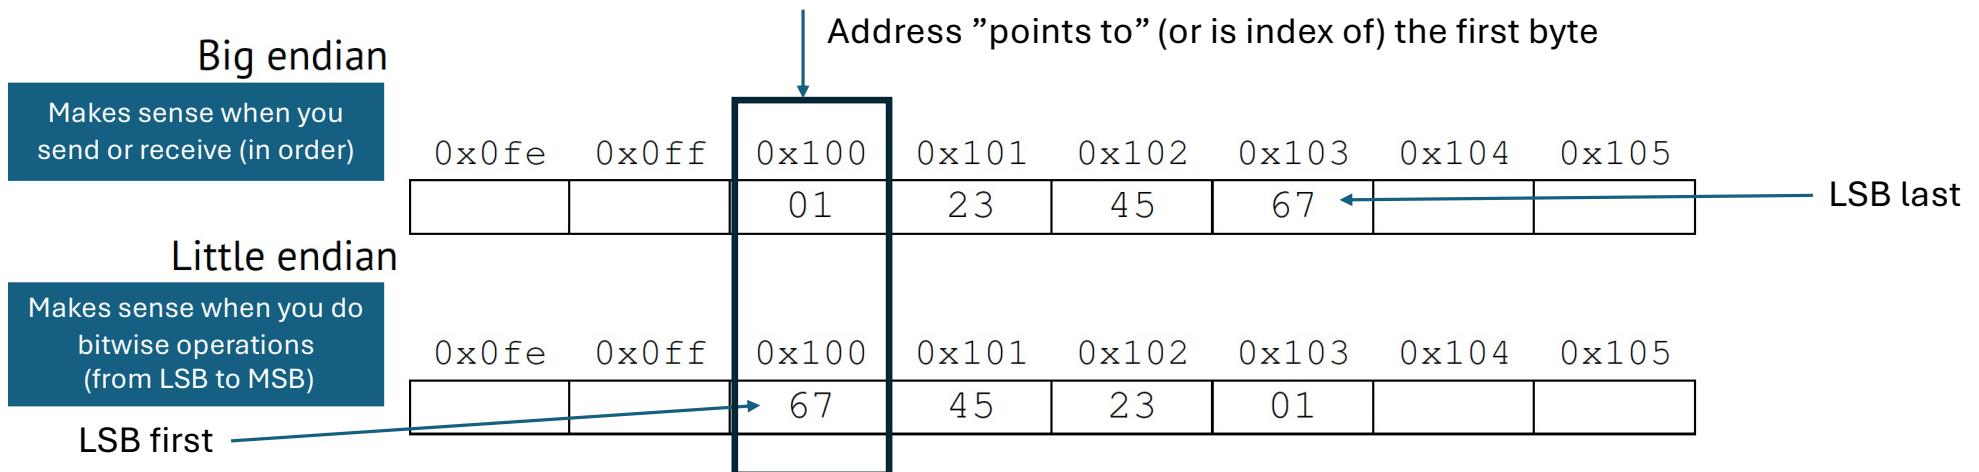
Example

- Value encoded on 4 bytes (e.g., `uint32_t`):

$X = 0x01234567$

LSB

- *Address of this value &X = 0x100*



Examining data representation

- **Code to print byte representation of data**

- ▶ Casting pointer to `unsigned char*` allows treatment as byte array.

```
void show_bytes(unsigned char* start, size_t len) {  
    size_t i;  
    for (i = 0; i < len; i++)  
        printf("%p\t0x%.2x\n", start+i, start[i]);  
    printf("\n");  
}
```

printf directives:

- `%p`: Print pointer.
- `%x`: Print hexadecimal.

Text representation

A text is a sequence of characters.

How to interpret binary numbers into characters?

Character sets

A character set maps a *number* to a *character*.

- **ASCII** maps characters to the range 0—127 (asciitable.com).
 - A character (char) is represented on 8 bits.
 - Some are invisible/unprintable control characters
- Unicode is a superset of ASCII that defines tens of thousands of characters for all the world's scripts.
 - UTF8, UTF16, UTF32 encodings on up to 4 bytes

ASCII (mapping) table

Control characters			Normal characters														
000	nul	016	dle	032	«	048	0	064	@	080	P	096	‘	112	p		
001	soh	017	dc1	033	!	049	1	065	A	081	Q	097	a	113	q		
002	stx	018	dc2	034	“	050	2	066	B	082	R	098	b	114	r		
003	etx	019	dc3	035	#	051	3	067	C	083	S	099	c	115	s		
004	eot	020	dc4	036	\$	052	4	068	D	084	T	100	d	116	t		
005	enq	021	nak	037	%	053	5	069	E	085	U	101	e	117	u		
006	ack	022	syn	038	&	054	6	070	F	086	V	102	f	118	v		
007	bel	023	etb	039	‘	055	7	071	G	087	W	103	g	119	w		
008	bs	024	can	040	(056	8	072	H	088	X	104	h	120	x		
009	tab	025	em	041)	057	9	073	I	089	Y	105	i	121	y		
010	lf	026	eof	042	*	058	:	074	J	090	Z	106	j	122	z		
011	vt	027	esc	043	+	059	;	075	K	091	[107	k	123	{		
012	np	028	fs	044	‘	060	<	076	L	092	«	108	l	124			
013	cr	029	gs	045	-	061	=	077	M	093]	109	m	125	}		
014	so	030	rs	046	.	062	>	078	N	094	^	110	n	126	~		
015	si	031	us	047	/	063	?	079	O	095	_	111	o	127	del		

How printf works

```
int x = 1234;  
printf("x: %d\n", x);
```

The text *string* that is passed to `printf()` looks like this in memory:

Characters	x	:		%	d	\n	\0
Bytes	120	58	32	37	100	10	0

`printf()` rewrites format specifiers (`%d`) to the textual representation of their corresponding value argument:

Characters	x	:		1	2	3	4	\n	\0
Bytes	120	58	32	49	50	51	52	10	0

These bytes (except the `0`) are then written to *standard output* (typically the terminal) which interprets them as characters and eventually draws pixels on the screen.

How about files?

Binary file: sequence of bytes

Text file: sequence of characters (e.g., ASCII)

Writing bytes

```
#include <stdio.h>

int main() {
    // Open for writing ("w")
    FILE *f = fopen("output", "w");

    char c = 42;

    fwrite(&c, sizeof(char), 1, f);

    fclose(f);
}
```

- Produces a file output.
- File contains the byte 42, corresponding to the ASCII character *.
- **char is just an 8-bit integer type!**
 - ▶ No special “character” meaning.
 - ▶ Most Unicode characters will not fit in a single char (e.g. ’æ’ needs 16 bits in UTF-8).
 - ▶ Name is unfortunate/historical.
 - ▶ Signedness is *implementation-defined* for historical reasons.

Converting an non-negative integer to ASCII

```
FILE *f = fopen("output", "w");
int x = 1337;                      // Number to write;
char s[10];                         // Output buffer.
int i = 10;                          // Index of last character written.
while (1) {
    int d = x % 10;                  // Pick out last decimal digit.
    x = x / 10;                      // Remove last digit.
    i = i - 1;                        // Index of next character.
    s[i] = '0' + d;                  // Save ASCII character for digit.
    if (x == 0) { break; } // Stop if all digits written.
}
fwrite(&s[i], sizeof(char), 10-i, f); // Write ASCII bytes.
fclose(f);                           // Close output file.
```

Reading all bytes in a file

```
#include <stdio.h>
#include <assert.h>

int main(int argc, char* argv[]) {
    FILE *f = fopen(argv[1], "r");
    unsigned char c;
    while (fread(&c, sizeof(char), 1, f) == 1) {
        printf("%3d", (int)c);
        if (c > 31 && c < 127) {
            fwrite(&c, sizeof(char), 1, stdout);
        }
        printf("\n");
    }
}
```

Reading all bytes in a file

```
$ gcc -o fread-bytes -Wall -Wextra -pedantic fread-bytes.c
$ ./fread-bytes fread-bytes.c
35 #
105 i
110 n
99 c
108 l
117 u
100 d
101 e
32
60 <
...
...
```

How about files?

```
a.out+-----+  
1 00000000: cffa edfe 0c00 0001 0000 0000 0200 0000 .....  
2 00000010: 1100 0000 2004 0000 8500 2000 0000 0000 .....  
3 00000020: 1900 0000 4800 0000 5f5f 5041 4745 5a45 .....H.....PAGEZE  
4 00000030: 524f 0000 0000 0000 0000 0000 0000 0000 R0.....  
5 00000040: 0000 0000 0100 0000 0000 0000 0000 0000 .....  
6 00000050: 0000 0000 0000 0000 0000 0000 0000 0000 .....  
7 00000060: 0000 0000 0000 0000 1900 0000 8801 0000 .....  
8 00000070: 5f5f 5445 5854 0000 0000 0000 0000 0000 _TEXT.....  
9 00000080: 0000 0000 0100 0000 0040 0000 0000 0000 .....@.....  
10 00000090: 0000 0000 0000 0000 0040 0000 0000 0000 .....@.....  
11 000000a0: 0500 0000 0500 0000 0400 0000 0000 0000 .....  
12 000000b0: 5f5f 7465 7874 0000 0000 0000 0000 0000 _text.....  
13 000000c0: 5f5f 5445 5854 0000 0000 0000 0000 0000 _TEXT.....  
14 000000d0: 6004 0000 0100 0000 f000 0000 0000 0000 .....  
15 000000e0: 6004 0000 0200 0000 0000 0000 0000 0000 .....  
16 000000f0: 0004 0080 0000 0000 0000 0000 0000 0000 .....  
17 00000100: 5f5f 7374 7562 7300 0000 0000 0000 0000 _stubs.....  
18 00000110: 5f5f 5445 5854 0000 0000 0000 0000 0000 _TEXT.....  
19 00000120: 5005 0000 0100 0000 0c00 0000 0000 0000 P.....  
20 00000130: 5005 0000 0020 0000 0000 0000 0000 0000 P.....  
21 00000140: 0804 0080 0000 0000 0c00 0000 0000 0000 .....  
22 00000150: 5f5f 6373 7472 696e 6700 0000 0000 0000 _cstring.....  
23 00000160: 5f5f 5445 5854 0000 0000 0000 0000 0000 _TEXT.....  
24 00000170: 5c05 0000 0100 0000 0a00 0000 0000 0000 .....  
25 00000180: 5c05 0000 0000 0000 0000 0000 0000 0000 .....  
26 00000190: 0200 0000 0000 0000 0000 0000 0000 0000 .....  
27 000001a0: 5f5f 7562 7769 6e64 5f69 6e66 6f00 0000 _ unwind_info.....  
28 000001b0: 5f5f 5445 5854 0000 0000 0000 0000 0000 _TEXT.....  
29 000001c0: 6805 0000 0100 0000 5800 0000 0000 0000 h.....X.....  
30 000001d0: 6805 0000 0020 0000 0000 0000 0000 0000 h.....  
31 000001e0: 0000 0000 0000 0000 0000 0000 0000 0000 .....  
.....
```

Compactness of storage

- A 32-bit integer takes up to 12 bytes to store as base-10 ASCII digits
 - 4 bytes as raw data
 - **Raw data takes up less space and is much faster to read.**
 - But we need special programs to decode the data to human-readable form.

Take-aways

- Memory is abstracted as an array of bytes.
- The unit of read and write to memory is a word (8B on 64 bits machines)
 - Little endian: LSB first (on all modern processors); big endian on network.
- Pointers are variables that contain the addresses of variables, i.e., indexes in the memory array.
 - As variable, a pointer can be stored in memory (we can thus have pointers to pointers ...)
- Text is just another interpretation of binary data (through character maps, e.g., ASCII or Unicode)
 - Interpreted when printing to terminal or when writing to/reading from a file (binary files are more compact than text files)